

▼ Interview Test

In this notebook a neural network is trained to learn to detect different classes of flowers.

Instruction to reproduce this notebook:

- download the compressed dataset file from this link: <https://www.robots.ox.ac.uk/~vgg/data/flowers/17/17flowers.tgz>
- copy the file in your Google Drive inside a directory called BNS_TEST (see the variable 'data_path' for the exact path)
- uncomment the cell where the zip directories are exacted.

The dataset is composed of 17 different classes of flowers. For each classes you can find 80 pictures. This is a balanced dataset. The images are ordered by class.

Example of image name: 'image_0828.jpg'.

Follows the list of the classes names with the corresponding index range (the number in the name) :

0 Daffodill 1/80
1 Snowdrop 81/160
2 Lilyvalley 161/240
3 Bluebell 241/320
4 Crocus 321/400
5 Iris 401/480
6 Tigerlily 481/560
7 Tulip 561/640
8 Fritillary 641/720
9 Sunflower 721/800
10 Daisy 801/880
11 Coltsfoot 881/960
12 Dandelion 961/1040
13 Cowslip 1041/1120
14 Buttercup 1121/1200
15 Windflower 1201/1280
16 Pansy 1281/1360

```
#mount your drive directory to access the data
from google.colab import drive
drive.mount('/content/drive/')
```

Drive already mounted at /content/drive/; to attempt to forcibly remount, call drive.mount("/content/drive/", force_remount=True).

```
from sklearn import metrics
import numpy as np
from matplotlib import pyplot as plt
import os
from warnings import filterwarnings
import tensorflow as tf
from tensorflow import io
from tensorflow import image
from tensorflow import keras
from tensorflow.keras import models
from tensorflow.keras import layers
import tarfile
from sklearn.utils import shuffle
from numpy.lib.function_base import extract
from sklearn.model_selection import train_test_split
```

```
tf.get_logger().setLevel('ERROR')
```

```
#path variables, check if is different in your system
data_path = '/content/drive/MyDrive/BNS_TEST/'
zip_path = data_path + '17flowers.tgz'
dataset_dir = data_path + 'dataset'
image_path= dataset_dir + '/jpg/image_0828.jpg'
```

```
img_height = 180
img_width = 180
BATCHSIZE = 32
```

```
# check hardware acceleration
device_name = tf.test.gpu_device_name()
print('GPU: ', device_name)
```

```
GPU:  /device:GPU:0
```

```
#!UNCOMMENT THIS CELL IF THE DATASET WAS NOT EXTRACTED ALREADY!
```


```
#extract the dataset zip directory in the drive
```

```
# open file
#file = tarfile.open(zip_path)
# extracting file
#file.extractall(dataset_dir)
#file.close()
```

```
#show image example
```

```
filterwarnings("ignore")
tf_img = io.read_file(image_path)
tf_img = image.decode_png(tf_img, channels=3)
print(np.amax(tf_img))
plt.imshow(tf_img)
tf_img.shape
```

```
255
TensorShape([500, 538, 3])
```



```
#create class names list
class_names= ['Daffodil', 'Snowdrop', 'LilyValley', 'Bluebell', 'Crocus', 'Iris', 'Tigerlily', 'Tulip', 'Fritillary', 'Sunflower', 'Daisy', 'Colts']
class_size = len(class_names)
```

```
#create labels list
# x=[0]*N list of size N, all N elements = 0.
labels = []
for i in range(17):
    labels += [i]*80
```

```
dataset_size=len(labels)
```

```
#create image paths list
filenames_path = dataset_dir + '/jpg/files.txt'
filenames = []
file = open(filenames_path, "r")
filenames = file.read().splitlines() #each line without \n
file.close()
filenames = [ dataset_dir + '/jpg/' + s for s in filenames]
```

```
#shuffle in a consistent way the 2 lists
filenames,labels = shuffle(filenames, labels, random_state=0)
print(filenames[0], labels[0])
```

```
#extract test
x_train, x_test, y_train, y_test = train_test_split(filenames, labels, test_size=0.15, random_state=4, stratify=labels )
len(x_train),len(x_test)
```

```
#parse every image in the dataset using `map`
#retrieve the image, convert it to a tensor and resize it to fit the input layer of the model 180x180
```

```

def _parse_function(filename, label):
    image_string = tf.io.read_file(filename)
    image_decoded = tf.image.decode_jpeg(image_string, channels=3)
    image = tf.cast(image_decoded, tf.float32)
    image = tf.image.resize(image, [img_height, img_width])
    return image, label

#test dataset
filenames_test = tf.constant(x_test)
labels_test = tf.constant(y_test)

dataset_test = tf.data.Dataset.from_tensor_slices((filenames_test, labels_test))

dataset_test = dataset_test.map(_parse_function)
print(tf.data.experimental.cardinality(dataset_test).numpy())

    204

#I dont use tf.keras.utils.image_dataset_from_directory (classes organized by directories needed), I create a tf.data.Dataset from a dire

filenames = tf.constant(x_train)
labels = tf.constant(y_train)

dataset = tf.data.Dataset.from_tensor_slices((filenames, labels))

dataset = dataset.map(_parse_function)

dataset_size=tf.data.experimental.cardinality(dataset).numpy()

#create train test and validation set.
val_size = int((dataset_size) * 0.2)
train_ds = dataset.skip(val_size)
val_ds = train_ds.take(val_size)

print(tf.data.experimental.cardinality(train_ds).numpy())
print(tf.data.experimental.cardinality(val_ds).numpy())

    925
    231

def configure_for_performance(ds):
    ds = ds.cache()
    ds = ds.shuffle(buffer_size=1000)
    ds = ds.batch(BATCHSIZE)
    ds = ds.prefetch(buffer_size=tf.data.AUTOTUNE) #allows later elements to be prepared while the current element is being processed
    return ds

train_ds = configure_for_performance(train_ds)
val_ds = configure_for_performance(val_ds)

```

The following is the architecture of the network, trained from scratch.

Since I was facing severe overfitting I added data augmentation layers, and a dropout layer to increase the performances.

```

model = models.Sequential()
model.add(tf.keras.Input(shape=(180,180,3)))

#Normalization layer
model.add(tf.keras.layers.Rescaling(1./255)),

#Data augmentation layers for overfitting
model.add(layers.experimental.preprocessing.RandomFlip(mode='horizontal', seed=123))
model.add(layers.experimental.preprocessing.RandomRotation(factor=0.25, seed=123, fill_mode='nearest')) #90degree

model.add(layers.Conv2D(32,(3,3),activation = 'relu' )) #32 filters 3x3, 3 channels #num par= 3x3x3x32+32=896
model.add(layers.MaxPooling2D((2,2))) #reduce activation map size by +-half, every 4 entries take 1 (2columns,2rows)
model.add(layers.Conv2D(64, (3,3), activation = 'relu' )) #64 filters 3x3, 32 channels #num par 3x3x32x64+64= 18496
model.add(layers.MaxPooling2D(2,2))
model.add(layers.Conv2D(128,(3,3), activation='relu' )) #128 filters 3x3, 64 channels #num par 3x3x64x128+128=73856
model.add(layers.MaxPooling2D(2,2))
model.add(layers.Conv2D(128,(3,3), activation='relu' )) #128 filters 3x3, 128 channels #num par 3x3x128x128+128=147584
model.add(layers.MaxPooling2D(2,2))
model.add(layers.Flatten())

#overfitting
model.add(layers.Dropout(0.4))

model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(17,activation='softmax'))
model.summary()

```

```
#opt = keras.optimizers.RMSprop(learning_rate=0.001)
opt = keras.optimizers.Adam()

model.compile(
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    optimizer=opt,
    metrics=['accuracy']
)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 180, 180, 3)	0
random_flip (RandomFlip)	(None, 180, 180, 3)	0
random_rotation (RandomRotation)	(None, 180, 180, 3)	0
conv2d (Conv2D)	(None, 178, 178, 32)	896
max_pooling2d (MaxPooling2D)	(None, 89, 89, 32)	0
conv2d_1 (Conv2D)	(None, 87, 87, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 43, 43, 64)	0
conv2d_2 (Conv2D)	(None, 41, 41, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 20, 20, 128)	0
conv2d_3 (Conv2D)	(None, 18, 18, 128)	147584
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 128)	0
flatten (Flatten)	(None, 10368)	0
dropout (Dropout)	(None, 10368)	0
dense (Dense)	(None, 512)	5308928
dense_1 (Dense)	(None, 17)	8721
Total params: 5,558,481		
Trainable params: 5,558,481		
Non-trainable params: 0		

```
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=24
)
```

```
Epoch 1/24
29/29 [=====] - 24s 206ms/step - loss: 2.6236 - accuracy: 0.1124 - val_loss: 2.2546 - val_accuracy: 0.2468
Epoch 2/24
29/29 [=====] - 3s 104ms/step - loss: 2.1277 - accuracy: 0.2659 - val_loss: 1.9092 - val_accuracy: 0.3117
Epoch 3/24
29/29 [=====] - 3s 96ms/step - loss: 1.8807 - accuracy: 0.3578 - val_loss: 1.6646 - val_accuracy: 0.4156
Epoch 4/24
29/29 [=====] - 3s 101ms/step - loss: 1.6364 - accuracy: 0.4227 - val_loss: 1.3428 - val_accuracy: 0.5671
Epoch 5/24
29/29 [=====] - 4s 131ms/step - loss: 1.4266 - accuracy: 0.5038 - val_loss: 1.3245 - val_accuracy: 0.5411
Epoch 6/24
29/29 [=====] - 3s 100ms/step - loss: 1.2909 - accuracy: 0.5643 - val_loss: 1.1398 - val_accuracy: 0.5887
Epoch 7/24
29/29 [=====] - 3s 96ms/step - loss: 1.2801 - accuracy: 0.5654 - val_loss: 1.1780 - val_accuracy: 0.6320
Epoch 8/24
29/29 [=====] - 4s 139ms/step - loss: 1.1509 - accuracy: 0.6130 - val_loss: 1.0103 - val_accuracy: 0.6537
Epoch 9/24
29/29 [=====] - 3s 104ms/step - loss: 1.1456 - accuracy: 0.6205 - val_loss: 1.0026 - val_accuracy: 0.6753
Epoch 10/24
29/29 [=====] - 3s 97ms/step - loss: 0.9508 - accuracy: 0.6692 - val_loss: 0.8510 - val_accuracy: 0.6840
Epoch 11/24
29/29 [=====] - 3s 98ms/step - loss: 0.9251 - accuracy: 0.6962 - val_loss: 0.7783 - val_accuracy: 0.7143
Epoch 12/24
29/29 [=====] - 3s 118ms/step - loss: 0.8393 - accuracy: 0.7124 - val_loss: 0.7504 - val_accuracy: 0.7532
Epoch 13/24
29/29 [=====] - 4s 120ms/step - loss: 0.7744 - accuracy: 0.7286 - val_loss: 0.6057 - val_accuracy: 0.7965
Epoch 14/24
```

```

29/29 [=====] - 3s 95ms/step - loss: 0.6654 - accuracy: 0.7838 - val_loss: 0.5945 - val_accuracy: 0.8139
Epoch 15/24
29/29 [=====] - 3s 93ms/step - loss: 0.6478 - accuracy: 0.7859 - val_loss: 0.4934 - val_accuracy: 0.8312
Epoch 16/24
29/29 [=====] - 3s 96ms/step - loss: 0.6087 - accuracy: 0.7849 - val_loss: 0.5574 - val_accuracy: 0.8009
Epoch 17/24
29/29 [=====] - 3s 100ms/step - loss: 0.6657 - accuracy: 0.7751 - val_loss: 0.5475 - val_accuracy: 0.8139
Epoch 18/24
29/29 [=====] - 3s 103ms/step - loss: 0.5861 - accuracy: 0.8022 - val_loss: 0.5058 - val_accuracy: 0.8312
Epoch 19/24
29/29 [=====] - 3s 120ms/step - loss: 0.5147 - accuracy: 0.8270 - val_loss: 0.4982 - val_accuracy: 0.8225
Epoch 20/24
29/29 [=====] - 3s 120ms/step - loss: 0.5791 - accuracy: 0.8054 - val_loss: 0.4813 - val_accuracy: 0.8485
Epoch 21/24
29/29 [=====] - 3s 95ms/step - loss: 0.5088 - accuracy: 0.8303 - val_loss: 0.3877 - val_accuracy: 0.8615
Epoch 22/24
29/29 [=====] - 3s 95ms/step - loss: 0.4755 - accuracy: 0.8346 - val_loss: 0.3693 - val_accuracy: 0.8615
Epoch 23/24
29/29 [=====] - 5s 163ms/step - loss: 0.4595 - accuracy: 0.8443 - val_loss: 0.2734 - val_accuracy: 0.9091
Epoch 24/24
29/29 [=====] - 3s 99ms/step - loss: 0.3955 - accuracy: 0.8692 - val_loss: 0.2398 - val_accuracy: 0.9221

```

```
#print the plot of the loss and accuracy values during the various epochs of the training
```

```
import matplotlib.pyplot as plt
```

```
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
```


```
epochs = range(len(acc))
```

```
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
```

```
plt.figure()
```

```
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
```

```
plt.show()
```



```
#example of single inference using a personal picture. save the pic in the Drive directory and add the filename in the following line
sunflower_path = data_path + 'pic.jpg'
```

```
img = tf.keras.utils.load_img(
    sunflower_path, target_size=(img_height, img_width)
)
```

```

img_array = tf.keras.utils.img_to_array(img)
img_array = tf.expand_dims(img_array, 0) # Create a batch

predictions = model.predict(img_array)
print(max(predictions[0]))

score = tf.nn.softmax(predictions[0])
print(score)
print(sum(score))
print(
    "This image most likely belongs to {} with a {:.2f} percent confidence."
    .format(class_names[np.argmax(score)], 100 * np.max(score))
)

1/1 [=====] - 0s 252ms/step
1.0
tf.Tensor(
[0.05342371 0.05342371 0.05342371 0.05342371 0.05342371 0.05342371
 0.05342371 0.05342371 0.05342371 0.14522068 0.05342371 0.05342371
 0.05342371 0.05342371 0.05342371 0.05342371 0.05342371], shape=(17,), dtype=float32)
tf.Tensor(0.99999994, shape=(), dtype=float32)
This image most likely belongs to Sunflower with a 14.52 percent confidence.

```

#nfer the test images using the trained model

```

dataset_test = dataset_test.batch(BATCHSIZE)
test_pred = model.predict(dataset_test)
predicted_labels = np.argmax(test_pred, axis=1)

test_labels = np.array(y_test)
test_predictions = np.squeeze(predicted_labels)

```

```
7/7 [=====] - 1s 117ms/step
```


```
m = metrics.confusion_matrix(test_labels, test_predictions)
```

```

print( f'report:\n {metrics.classification_report(test_labels, test_predictions)}')
plt.matshow(m, cmap=plt.cm.get_cmap('Blues', 16))
plt.colorbar()
plt.show()

```

report:	precision	recall	f1-score	support
0	0.36	0.42	0.38	12
1	0.82	0.75	0.78	12
2	1.00	0.75	0.86	12
3	0.73	0.67	0.70	12
4	0.75	0.75	0.75	12
5	0.90	0.75	0.82	12
6	0.85	0.92	0.88	12
7	0.33	0.25	0.29	12
8	0.92	0.92	0.92	12
9	0.80	1.00	0.89	12
10	0.92	0.92	0.92	12
11	0.44	0.67	0.53	12
12	0.57	0.67	0.62	12
13	0.62	0.67	0.64	12
14	0.57	0.33	0.42	12
15	1.00	0.83	0.91	12
16	0.79	0.92	0.85	12
accuracy			0.72	204
macro avg	0.73	0.72	0.71	204
weighted avg	0.73	0.72	0.71	204



✓ 0 s data/ora di completamento: 23:16

